

# CS 556 Mathematical Foundations of Machine Learning

Linear Regression

# Today's Lecture

- Linear regression
- Optimization
- Features
- Beyond linear regression models
- Bias and variance

# In Praise of Linear Models!

- Despite its simplicity, the linear model has distinct advantages in terms of its **interpretability** and often shows good **predictive performance**.
- In this lecture, we discuss (1) optimization, (2) some ways in which the simple linear model can be improved, by replacing ordinary least squares fitting with some alternative fitting procedures.

# Regression Task

- Example: given the gender and height information, predict the weight.

Data sample	Gender	Height (inches)	Y: Weight (pounds)
1	Male	73	190
2	Male	67	165
3	Female	65	130

Features/Attributes (X)

Output (Y)




$$y \approx f(x) = \sum_{i=1}^d w_i x_i = xw^T$$

# Linear Regression

- **Input:**  $d$ -dimensional  $x = (x_1, x_2, \dots, x_d)$ .
- **Output:** Real value output  $y$ .
- **Regression Model:**

$$y \approx f(x) = \sum_{i=1}^d w_i x_i = x w^T$$



where  $w = (w_1, w_2, \dots, w_d)$  are parameters (weights/coefficients).

- **Training Data:**  $(X_1, Y_1), (X_2, Y_2), \dots, (X_N, Y_N)$  used for the estimation of the parameters  $w$ .

# The Learning Problem

- **Hypothesis class:** Consider some restricted set  $F$  of mappings  $f: X \rightarrow Y$  from input data to output.
- **Estimation:** Find the best estimate of mapping  $\hat{f} \in F$  based on the training data.
- **Evaluation:** Measure how well  $\hat{f}$  generalizes to the testing data (unseen examples), i.e., whether  $\hat{f}(x_{new})$  agrees with  $y_{new}$ .

# Estimation Criterion

- **Loss function**  $Loss(x, y, w)$ : a function of  $w$  that quantifies how well (more accurately, how badly) you would be if you use  $w$  to make a prediction on  $x$  when the correct output is  $y$ .
- **Objective**: the estimation problem can be solved by minimizing the loss function.
- Valid for any parameterized class of mapping from examples to predictions.
- Valid when the predictions are discrete labels or real valued as long as the loss is defined properly.

# Objective of Linear Regression

- **Prediction**: for one data example  $(x, y)$ , the prediction is:

$$f_w(x) = xw^T$$

- **Residual**: difference between the prediction and the target.
- **Squared loss**:  $Loss(x, y, w) = (y - f_w(x))^2$ ; smaller values of loss indicates that  $f_w(x)$  is a good approximation of  $y$ .
- **Example**:
  - $w = [2, -1], x = [2, 0], y = -1, L(x, y, w) = 25$



# Objective of Linear Regression

- Loss is easy to minimize if there is only one example.
- **Overall training loss**: on the training data  $D_{train}$  with  $N$  examples:

$$\min_{w \in R^d} TrainLoss(w) = \min_{w \in R^d} \frac{1}{N} \sum_{(x,y) \in D_{train}} L(x, y, w)$$


- **Key**: learn the parameters to make the global tradeoffs.
- If  $Loss(x, y, w) = (y - f_w(x))^2$ , the residual sum of squares (RSS) is defined as  $RSS = \sum_{(x,y) \in D_{train}} L(x, y, w)$ .

# Regression Loss Functions

- Squared loss:  $L_2(x, y, w) = (y - f_w(x))^2$ 
  - The overall loss will be mean squared error (MSE)
  - Root mean squared error (RMSE):  $RMSE = \sqrt{MSE}$
- Absolute loss:  $L_1(x, y, w) = |y - f_w(x)|$ 
  - The overall loss will be mean absolute error (MAE)

# How to Optimize: Gradient Descent

- A first-order iterative optimization algorithm for finding the local minimum of a differentiable function.
- To get the parameter that minimizes the objective function, we **iteratively move in the opposite direction of the gradient of the function**.
- The size of each step is determined by a parameter which is called **Learning Rate**.



# How to Optimize: Gradient Descent

- Objective:  $TrainLoss(w) = \min_{w \in \mathbb{R}^d} \frac{1}{N} \sum_{(x,y) \in D_{train}} (y - xw^T)^2$
- Gradient with respect to  $w$ : the direction that increases the loss the most.

$$\nabla_w TrainLoss(w) = \frac{1}{N} \sum_{(x,y) \in D_{train}} 2(y - xw^T) x$$

- Gradient descent update:

$$w \leftarrow w - \boxed{\eta} * \boxed{\nabla_w TrainLoss(w)}$$

Step size

Gradient

# Gradient Descent

---

```
1: procedure BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:      $g^{(i)}(w) = \text{evaluate\_gradient}(\text{TrainLoss}, \text{data}, w)$ 
4:      $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```

---

## Some Limitations:

- Each iteration requires going over **all training examples (Batch gradient descent)** – expensive and slow when have lots of data.
- Intractable for datasets that don't fit in memory.
- Guaranteed to converge to the global minimum for convex functions, but may end up at a local minimum for non-convex functions.

# Stochastic Gradient Descent (SGD)

- Considers **only one point at a time** to update weights.
- Not calculate the total error for whole data in one step, instead we calculate the error of each point and use it to update weights.

- Gradient descent update:

$$w \leftarrow w - \underbrace{\eta}_{\text{Step size}} * \underbrace{\nabla_w \text{TrainLoss}(w)}_{\text{Gradient}}$$

- Stochastic Gradient Descent:

For each  $(x, y) \in D_{\text{train}},$

$$w \leftarrow w - \eta \nabla_w \text{Loss}(x, y, w)$$

# Stochastic Gradient Descent (SGD)

---

```
1: procedure STOCHASTIC GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for example  $\in$  data do
5:        $g^{(i)}(w) = \text{evaluate\_gradient}(\text{loss}, \text{example}, w)$ 
6:        $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```

---

- **Shuffling**: to ensure that each data point creates an “independent” change on the model, without being biased by the same points before them.
- Easy to fit in memory as only one data point needs to be processed at a time, thus, computationally less expensive.
- With a high variance that cause the objective function to fluctuate heavily.
- May never reach local minima and oscillate around it due to the fluctuations in each step.

# Mini-batch Gradient Descent

---

```
1: procedure MINI-BATCH GRADIENT DESCENT
2:   for  $i$  in range(epochs) do
3:     np.random.shuffle(data)
4:     for batch  $\in$  get_batches(data, batch_size=50) do
5:        $g^{(i)}(w) = \text{evaluate\_gradient}(\text{loss}, \text{batch}, w)$ 
6:        $w = w - \text{learning\_rate} * g^{(i)}(w)$ 
```


---

- Instead of using the whole data for calculating gradient, we use only a mini-batch of it (batch size is a hyperparameter).
- Reduces the variance of the parameter updates, which can lead to more stable convergence.
- Can make use of highly optimized matrix optimizations common to state-of-the-art deep learning libraries that make computing the gradient w.r.t. a mini-batch very efficient.



# Learning Rate

Choosing an appropriate value of learning rate is important because it helps in determining how much we have to descent in each iteration.



# Summary so far

- Linear predictors:

$$f(x) = \sum_{i=1}^d w_i x_i = xw^T$$

- Loss minimization: learning as optimization

$$\min_{w \in \mathbb{R}^d} \text{TrainLoss}(w)$$

- Optimization algorithm: gradient decent or other variants

$$w \leftarrow w - \eta * \nabla_w \text{TrainLoss}(w)$$

# Today's Lecture

- Linear regression
- Optimization
- **Features**
- Beyond linear regression models
- Bias and variance

# Input Features

➤ The input features  $x_i$  can come from different sources:

- Quantitative inputs
- Transformations of quantitative inputs, such as log and square-root.
- Basis expansions, such as  $x_2 = x_1^2$  leading to a polynomial representation.
- Numeric or 'dummy' coding of the levels of qualitative inputs.
  - If  $G$  is a five-level factor input, we can create  $x_i, i = 1, \dots, 5$  such that  $x_i = I(G = i)$ .
  - This group of features represent the effect of  $G$  by a set of level-dependent constraints, since only one  $x_i$  is 1, and the others are 0.
- Interactions between features, for example:  $x_3 = x_1 x_2$ .

# Feature Extraction

- Example task: predict the weight of a person
- Question: what properties of a person might be relevant for predicting the weight  $y$ ?
- Feature extractor: given input  $x$ , output a set of (feature name, feature value) pairs.

<i>Gender:</i>	<i>Height:</i>	<i>Age:</i>
<i>Female</i>	<i>165 (cm)</i>	<i>25</i>

# Feature Vector Representation

- Mathematically, feature vector doesn't need feature names.

$$\begin{bmatrix} \textit{Gender:} & \textit{Height:} & \textit{Age:} \\ \textit{Female} & 165 \text{ (cm)} & 25 \end{bmatrix} \rightarrow [0 \quad 165 \quad 25]$$

- Feature vector representation: For input  $x$ ,  $x$  is a point in the high-dimensional space :

- Array representation: good for dense features

$$x = [x_1, x_2, \dots, x_d]$$

- Map representation: good for sparse features

$$x = \{\textit{Gender: 0, Height: 165, Age: 25}\}$$

# Linear Predictors

Weight vector  $w \in R^d$

$$\begin{bmatrix} \textit{Gender} & \textit{height} & \textit{age} \\ 1.5 & 0.8 & 1.1 \end{bmatrix}$$

Feature vector  $x \in R^d$

$$\begin{bmatrix} \textit{Gender} & \textit{height} & \textit{age} \\ 0 & 165 & 25 \end{bmatrix}$$

Predicting score:


$$f(x) = \sum_{i=1}^d w_i x_i = xw^T$$

# Feature Scaling

- Data may include features with **varying magnitudes**.

$\begin{bmatrix} \text{Gender:} & \text{Height:} & \text{Age:} \\ \text{Female} & 165 \text{ (cm)} & 25 \end{bmatrix}$

- ML algorithms may assign higher weight to greater values regardless of the unit of the values (e.g. 165 cm > 1.95 m), and hence provide wrong prediction.
- Normalization**: each feature contributes approximately proportionally.





# Feature Scaling

- Gradient descent converges much faster with feature scaling than without it.





Figure: <https://medium.com/analytics-vidhya/feature-scaling-normalization-standardization-and-scaling-c920ed3637e7>

# Feature Scaling: Standardization

- **Z-score normalization:** transform the data such that the resulting data has a mean of 0 and a standard deviation of 1.
- For each feature  $x_i$ , calculate the mean  $m_i$ , and the standard deviation  $\sigma_i$ .
- **Scaling features:**

$$x_i \leftarrow \frac{x_i - m_i}{\sigma_i}$$



# Feature Scaling: Min-Max Scaling

- Transform the data to  $[x_{min}^{new}, x_{max}^{new}]$ .


- Scaling features:

$$x_i \leftarrow \frac{x_i - x_{min}}{x_{max} - x_{min}} (x_{max}^{new} - x_{min}^{new}) + x_{min}^{new}$$

- If the new range is  $[0,1]$ ,

$$x_i \leftarrow \frac{x_i - x_{min}}{x_{max} - x_{min}}$$

- Example: If the weight of the students is in the range of  $[100, 150]$ , after normalization to  $[0,1]$ , the weight 120 is mapped to  $\frac{120-100}{150-100} = 0.4$



# Closed-form Solution

- Learning objective with  $(X_{N \times d}, Y_{N \times 1}, w_{d \times 1}^T)$

$$\min_{w \in \mathbb{R}^d} \frac{1}{N} \sum_{i=1}^N (Y_i - X_i w^T)^2 = \min (Y - Xw^T)^T (Y - Xw^T)$$

- Differentiating w.r.t  $w$ :

$$X^T (Y - Xw^T) = 0$$

- Setting the gradient to zero, we get closed form of estimates:

$$w = (X^T X)^{-1} X^T Y$$

- The resulting prediction errors  $\varepsilon_i = Y_i - X_i w^T$  are uncorrelated with any linear function of the inputs  $X$ .
- Easily extend to non-linear functions of the inputs.

# Beyond Linear Regression

- Prediction driven by score:  $f(x) = xw^T$
- Score is linear function of  $w$ , which permits efficient learning.
- Predictors can be expressive **non-linear** functions and decision boundaries of  $x$ .
- Example:  $m^{th}$  order polynomial regression on one-dimensional input where  $f: R \rightarrow R$  is given by:

$$f(w) = w_0 + w_1x + \cdots + w_{m-1}x^{m-1} + w_mx^m$$

# Polynomial Regression

- Polynomial regression is a **special case of linear regression**.

- Problem:

$$f(w) = w_0 + w_1x + \cdots + w_{m-1}x^{m-1} + w_mx^m$$

- Solution:


$$w = (X^T X)^{-1} X^T Y$$

- Where


$$w^T = \begin{pmatrix} w_0 \\ w_1 \\ \vdots \\ w_m \end{pmatrix} \quad X = \begin{pmatrix} 1 & x_1 & x_1^2 & \cdots & x_1^m \\ 1 & x_2 & x_2^2 & \cdots & x_2^m \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_N & x_N^2 & \cdots & x_N^m \end{pmatrix}$$

# Polynomial Regression


- Can be generally used when linear regression fails.
- Does not require the linear relationship between  $x$  and  $y$ .
- Predicts the best fit line that follows the pattern (curve) of the data.
- When the degree increases, the performance tend to increase, however, it also increases the risk of overfitting and underfitting.




degree = 1



degree = 3



degree = 5




degree = 7

# Underfitting and Overfitting

	Training Accuracy	Testing Accuracy	Possible Reason
Underfitting	0.24	0.25	Model is too simple: <ul style="list-style-type: none"><li>• Low dimensional</li><li>• Heavily regularized</li><li>• Bad assumption</li></ul>
Overfitting	0.95	0.27	Model is too complex: <ul style="list-style-type: none"><li>• High dimensional or non-parametric</li><li>• Weakly regularized</li><li>• Not enough data</li></ul>



# Underfitting and Overfitting



# Bias and Variance


- Assuming we are predicting  $Y$  based on the features  $X$ . There is a function  $Y = f(X) + \varepsilon$  where the error term is normally distributed with mean 0 and variance  $\sigma^2$ .
- The goal for us is to find a function  $\hat{f}(X)$  that approximates the true function  $f(X)$ .
- Expected squared prediction error at a point  $X$  will be:

$$E \left[ \left( Y - \hat{f}(X) \right)^2 \right] = \left( E[\hat{f}(X)] - f(X) \right)^2 + \left( E[\hat{f}(X)^2] - E^2[\hat{f}(X)] \right) + \sigma^2$$

$$= \text{Bias}^2 + \text{Variance} + \sigma^2$$

- Bias: Error from incorrect modeling assumption.
- Variance: Error from random noise.

# Bias and Variance Tradeoff



# Solution for Large Bias

- If your model cannot even fit the training examples, then you have large bias (underfitting)
- If you can fit the training data, but large error on testing data, then you probably have large variance (overfitting)
- Solution: redesign your model
  - Add more features as input
  - A more complex model

# Solution for Large Variance

- The least squares estimates often have low bias but large variance.
- It is possible to trade in a little bias for a larger reduction in variance leading to a smaller MSE than least square.
- Solution:
  - More data (very effective, but not always practical)
  - Feature subset selection
  - Regularized regression

*With a large number of features, we often would like to determine a smaller subset that exhibit the strongest effects.*

# Stepwise Selection

- Forward-stepwise Search
  - Start with no features
  - Greedily include the most relevant feature
  - Stop when selected the desired number of features
- Backward-stepwise Search
  - Start with all features
  - Greedily remove the least relevant feature
  - Stop when selected the desired number of features
- Inclusion/Removal criteria uses cross-validation

# Regularized Regression

- Subset selection, in which variables are either retained or discarded, is a discrete process. It often exhibits high variance, and so doesn't reduce the prediction error of the full model.
- Instead of directly minimize the MSE, the regularized regression usually take the form:

$$\min_{w \in R^d} \frac{1}{N} \sum_{(x,y) \in D_{train}} (y - xw^T)^2 + \lambda P(w)$$

Where  $P(w)$  is the penalized term on the smoothness or complexity of the function  $w$ .

# Regularization

- $L_1$  regularization: Lasso regression

$$L_1(X, Y, w) = \sum_{i=1}^N (f_w(X_i, w) - Y_i)^2 + \lambda |w|$$

- $L_2$  regularization: Ridge regression:

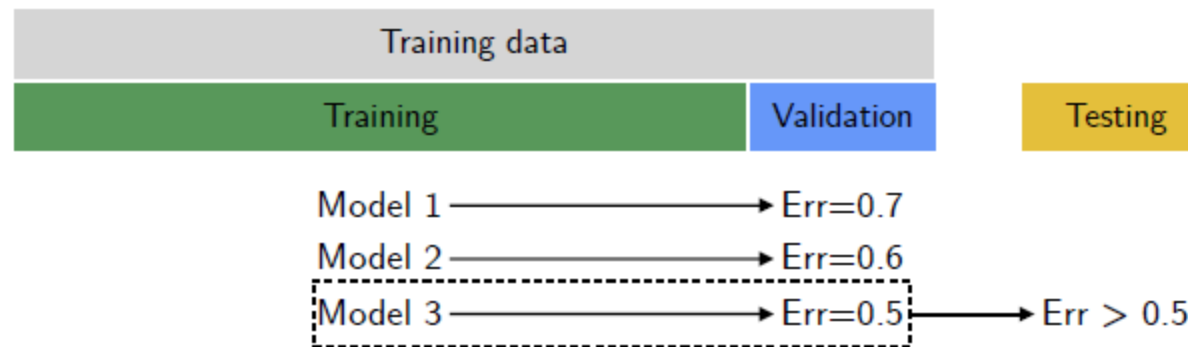
$$L_2(X, Y, w) = \sum_{i=1}^N (f_w(X_i, w) - Y_i)^2 + \lambda w^T w$$

- The tuning parameter  $\lambda$  serves to control the relative impact of these two terms on the regression coefficient estimates.
- Selecting a good value for  $\lambda$  is critical; cross-validation is used for this.



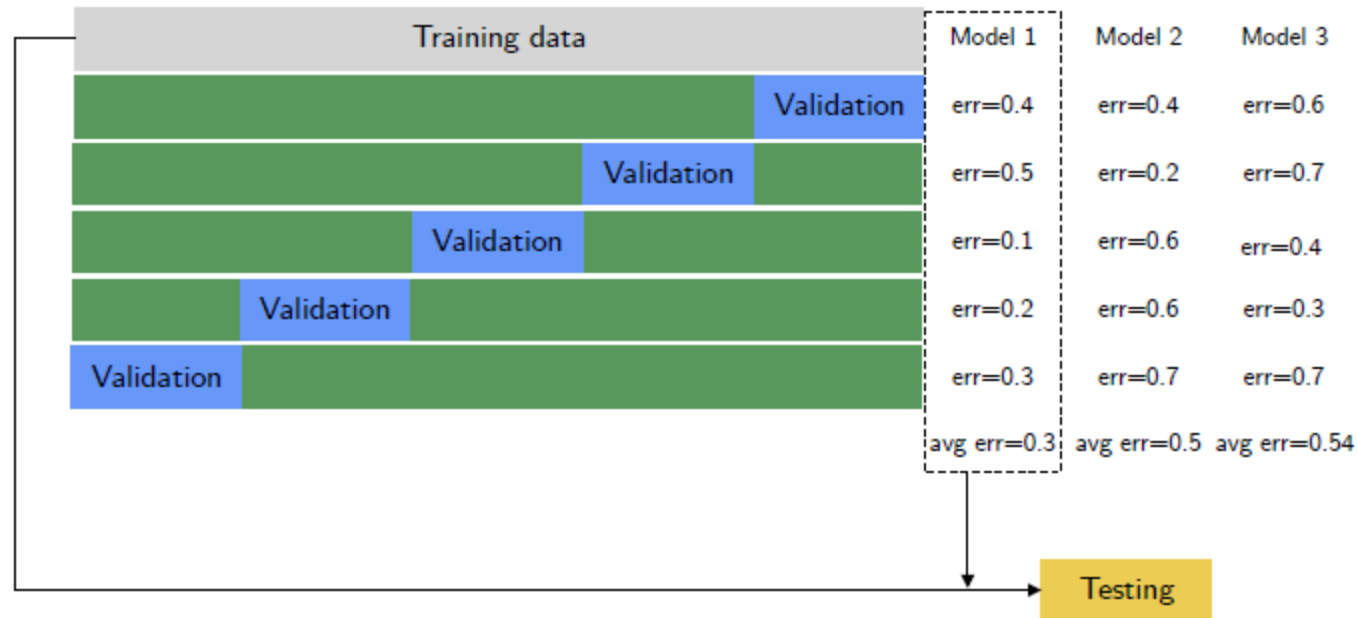
# Model Selection

- Usually a trade-off between bias and variance.
- We require a method to determine which of the models under consideration is best and can balance two kinds of error to minimize total error.
- Cross-validation: choose a grid of values for the parameters and compute the cross-validation error rate for each value.
- We then select the tuning parameter value for which the cross-validation error is smallest.



Training vs. Validation vs. Testing

# $k$ -fold Cross Validation



- In general: we perform  $k$  runs. Each run, it allows to use  $\frac{k-1}{k}$  of the available data for training.
- If the number of data is very limited, we can set  $k = N$  (total number of data points). This gives the leave-one-out cross-validation technique.

# Summary of Today's Lecture

- Linear regression
- Optimization
- Features
- Beyond linear regression models
- Bias and variance

# Slides

- Courtesy of Ping Wang.